OPPORTUNITY BRIEF · SALES AND PRESALES

Decolonized ethics and governance framework for generative AI in healthcare to prevent discrimination against vulnerable groups and ensure equitable access.

Healthcare × AI governance and assurance × Generative AI · OS027 v1 · refreshed 2026-08-18

ATTRACTIVENESS	RIGHT TO WIN	COMPETITION	SERVICEABLE MARKET	HORIZON	PORTFOLIO
60 is the world moving	**75** can we play, can we win	**HIGH** 0 named in the evidence	**€298m** partial evidence · per year	**LATER** research and standards activi...	**L0** 13 signals

Attractiveness and right to win are scored and displayed separately and are never combined (SC-12). Competition is a fourth quantity, not a discount on either.

The opportunity

This opportunity is about helping healthcare providers implement a decolonized ethics and governance framework for generative AI, ensuring that AI systems do not discriminate against vulnerable groups and that access to care remains equitable. It is for hospitals, health systems, and public health authorities that are deploying or planning to deploy generative AI and need to demonstrate trustworthy, human-centric AI governance. The need is urgent because AI in healthcare has been shown to amplify systemic inequalities, and regulators and payers are increasingly demanding accountability.

Market opportunity

Method	Total (TAM)	Serviceable (SAM)	Obtainable (SOM)	Confidence
Bottom-up: enterprises x adoption x contract value	**€484m** €39.7m – €2.3bn	**€298m** €24.4m – €1.4bn	**€17.9m** €1.5m – €85.2m	partial
Observed: tendered contract value, annualised	€70.8m	€70.8m	€4.3m	observed

All figures are annual, in EUR. Ranges compound the uncertainty of each factor below; the base case is the middle column of each.

Bottom-up: enterprises x adoption x contract value

Factor	Value	Basis	Source	As of
Enterprises in Healthcare (EU27_2020)	46,762 enterprises	observed	Eurostat · sbs_sc_ovw	2024
Enterprises adopting Generative AI	8.8 % of enterprises (10+ employees)	proxy	Eurostat · isoc_eb_ain2 · E_AI_TNLG	2025
Annual value of one engagement	€970k per enterprise per year	observed	TED (Tenders Electronic Daily) · ted	2026-05-19..2026-07-22
Size-weighted buyer base	5,698 enterprise-equivalents at large-enterprise engagement value	assumption	config/sizing.yaml · size_class_value_weights	size-2026-08-a
Assumed obtainable share of the serviceable market	6.0 % of SAM	assumption	config/sizing.yaml · obtainable_share	size-2026-08-a

Read this before quoting the number. Adoption rate is a declared proxy. Natural language generation Part of this vertical is not covered separately by the enterprise ICT survey, so the all-activities rate was used for those slices. Contract values are observed from EU public procurement (TED). For private-sector verticals they are used as a proxy for comparable enterprise deal size: the buying entity differs, the scope of work is comparable. Engagement value is scaled by enterprise size class (10-19: x0.04, 20-49: x0.09, 50-249: x0.35, GE250: x1.0), because the observed contract value comes from large-organisation tenders. Sector adoption is read from Eurostat's all-activities aggregate for part of this vertical, which the enterprise ICT survey does not cover separately: PROXY — ICT survey excludes health activities; PROXY — residential care

Observed: tendered contract value, annualised

Factor	Value	Basis	Source	As of
Tendered contract value in matching CPV categories	€70.8m per year	observed	TED (Tenders Electronic Daily) · ted	2026-05-19..2026-07-22
Assumed obtainable share of the serviceable market	6.0 % of SAM	assumption	config/sizing.yaml · obtainable_share	size-2026-08-a

Read this before quoting the number. Observed over 85 days of TED coverage and scaled to a year (x4.3); a short collection window makes that scaling rough. 7 matching notices, matched on vertical via the CPV crosswalk — §4.5.2's warning applies: a crosswalk error lands directly in this number. Public procurement only. Private-sector demand for the same capability is not visible here at all, so this is a floor rather than a market size.

The solution, in one picture

Decolonized AI Governance for Healthcare

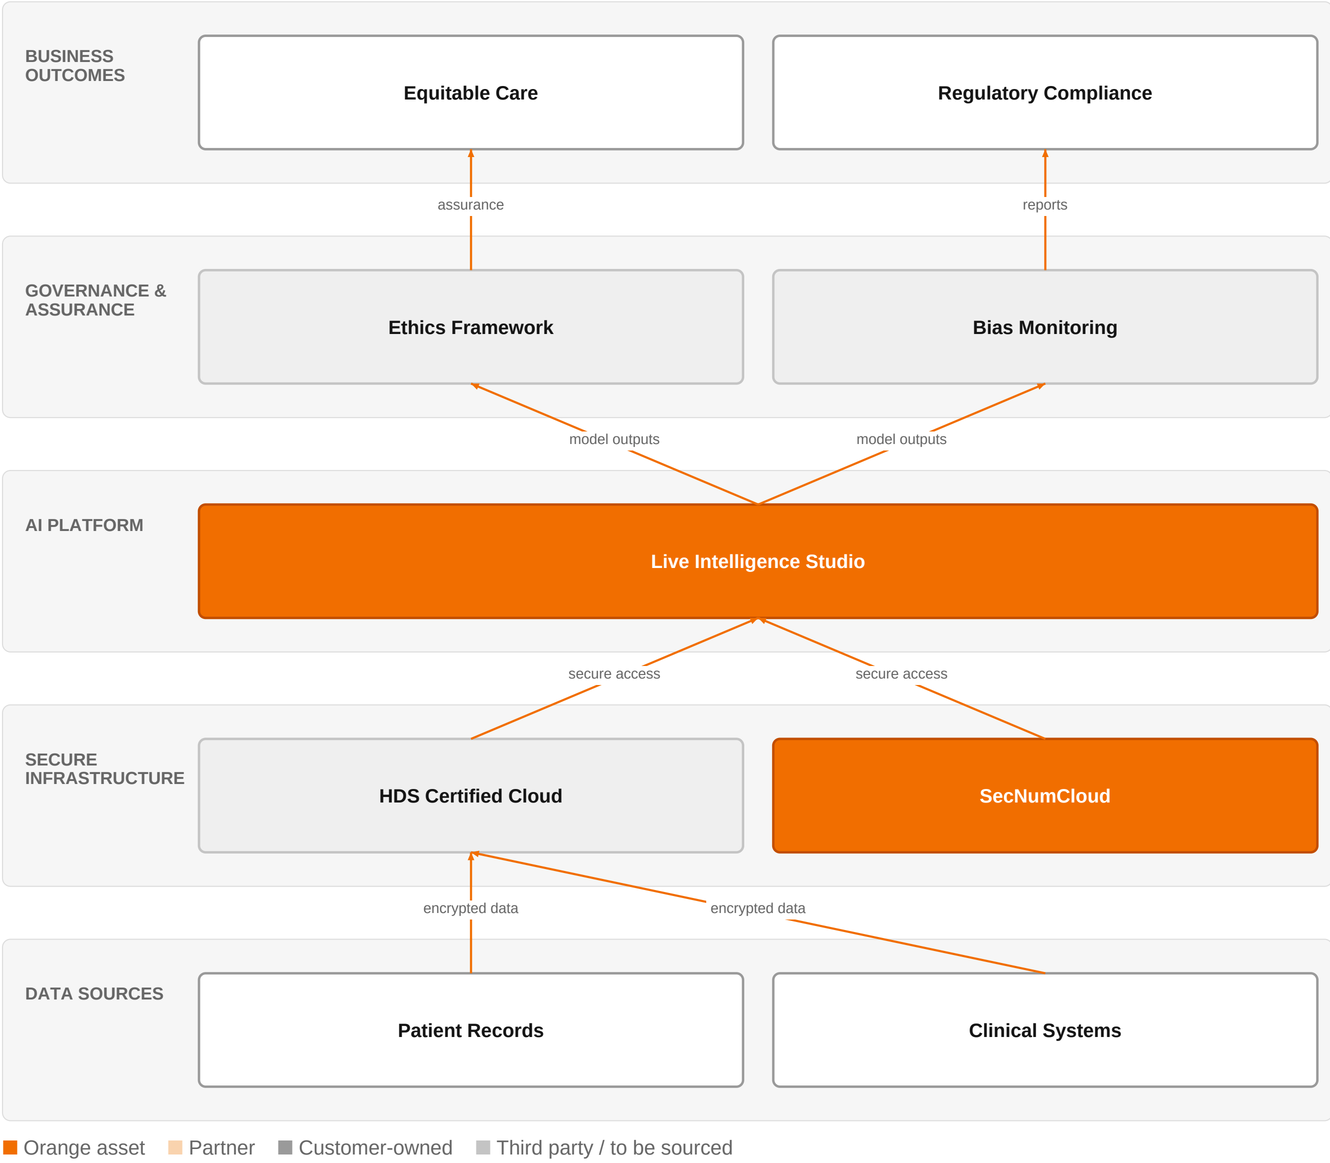

This picture shows how Orange's certified infrastructure and Live Intelligence Studio enable a governance layer that ensures generative AI in healthcare is equitable and compliant.

Boxes marked as Orange assets are validated against the linked business graph; a box the model claimed for Orange without a matching asset is shown as third party.

What has changed

Research and standards activity is highlighting that AI in healthcare often discriminates against vulnerable groups, amplifying systemic inequalities and preventing equitable access to care. This is driving a shift from merely 'trustworthy AI' to more comprehensive, decolonized frameworks that address structural biases. Additionally, the rapid digitalization of healthcare demands intelligent, low-latency, and privacy-preserving systems, and the rise of agentic AI and generative AI introduces new governance challenges. Hospitals are adopting AI for triage, imaging, and scheduling, but most deployments remain siloed, and a large proportion of healthcare AI pilots fail to scale due to governance issues.

`Sources: SIG-A4D6D84811FE cordis.europa.eu 2026-01-01; SIG-AECDD651DA98 openalex.org 2025-12-17; SIG-0A19EC698D37 openalex.org 2025-11-14; SIG-F2B27060ED4C arxiv.org 2026-08-06`

Who buys, and why now

The buyer is typically a Chief Medical Officer, Chief Information Officer, or Chief Data Officer at a large hospital or health system, or a director of health innovation in a public health authority. The pain is felt by clinical teams who see AI tools producing biased recommendations that disadvantage certain patient groups, and by compliance officers who fear regulatory and reputational risk. The trigger is often an adverse event, a regulatory audit, or a mandate from the board to ensure AI deployments are ethical and equitable.

`Sources: SIG-A4D6D84811FE cordis.europa.eu 2026-01-01`

Why it is hot now — every claim bound to a dated source

- AI has great potential in healthcare but raises ethical concerns as it often discriminates vulnerable groups, amplifying systemic inequalities. [SIG-A4D6D84811FE]

What Orange would deliver, and why it can win

What Orange would deliver

Orange would assemble a governance and assurance engagement using its Healthcare experts, AI/Data/Cloud experts, and Orange Group researchers to design a decolonized ethics framework tailored to the customer's generative AI use cases. The engagement would start with an assessment of existing AI deployments and data practices, then co-create a governance framework that includes bias detection and mitigation, transparency, and accountability mechanisms. Orange would leverage its Live Intelligence Studio to operationalize the framework, and its certifications (HDS, ISO 27001, SecNumCloud) to ensure sovereign, secure deployment. The delivery would be phased: first, a gap analysis; second, framework design; third, implementation support; and fourth, continuous monitoring and assurance.

Why Orange can win

Orange can win because it combines deep healthcare expertise with sovereign, certified infrastructure (HDS, SecNumCloud, ISO 27001) that is essential for handling sensitive health data. Its Live Intelligence Studio provides a platform to implement governance in practice, and its Orange Group researchers bring cutting-edge insights into trustworthy AI. However, the decolonized ethics framework itself is a nascent area, and Orange does not yet have a proven reference for this specific offering, so the win will depend on building credibility through thought leadership and partnerships.

Named assets behind this — each individually inspectable (LK-08)

Asset	Type	Link	Why it is linked	Curator
Live Intelligence Studio	offer	L0	offer addresses the use case AND provides the technology	unconfirmed
Microsoft	partner	L2	partner provides the required technology	unconfirmed
Google Cloud	partner	L2	partner provides the required technology	unconfirmed
NVIDIA	partner	L2	partner provides the required technology	unconfirmed
ISO 27001	certification	SUP	certification Orange holds is required by this vertical	unconfirmed
SecNumCloud	certification	SUP	certification Orange holds is required by this vertical	unconfirmed
HDS (Hebergeur de Donnees de Sante)	certification	SUP	certification Orange holds is required by this vertical	unconfirmed
AI / Data / Cloud experts	capability pool	SUP	capability pool staffs this domain, vertical or technology	unconfirmed
Cybersecurity experts	capability pool	SUP	capability pool staffs this domain, vertical or technology	unconfirmed
Defense and security experts	capability pool	SUP	capability pool staffs this domain, vertical or technology	unconfirmed
Healthcare experts	capability pool	SUP	capability pool staffs this domain, vertical or technology	unconfirmed
Orange Group researchers	capability pool	SUP	capability pool staffs this domain, vertical or technology	unconfirmed
MicroPort CardioFlow	reference	SUP	published reference in this vertical, different use case	unconfirmed
Bliksund	reference	SUP	published reference in this vertical, different use case	unconfirmed

Competitive landscape

HIGH competitive intensity (score 8.3; 0 of 8 listed players appear in this topic's own evidence)

Crowded, with heavyweight incumbents already visible in the evidence. Win on a specific differentiator or do not bid.

The competitive landscape is intense, with hyperscalers like Google Cloud and Microsoft offering generative AI platforms, and global system integrators like Accenture, Capgemini, IBM, Infosys, and TCS selling AI transformation services. These competitors are strong in technology and scale, but they may lack the sovereign, healthcare-specific certifications that Orange brings. IBM also competes in cybersecurity, which is relevant to AI governance. The customer will likely compare Orange's sovereign, certified approach against the broader AI platforms and consulting services of these players.

`Sources: SIG-702E3F720114 cordis.europa.eu 2025-11-01; SIG-F19E5C4832C5 cordis.europa.eu 2025-06-01`

Competitor	Category	Position here	Basis	Evidence
Google Cloud	Hyperscaler	sells Generative AI — also an Orange partner	structural	—
Microsoft	Hyperscaler	sells Generative AI — also an Orange partner	structural	—
Accenture	Global systems integrator	sells Generative AI	structural	—
Capgemini	Global systems integrator	sells Generative AI	structural	—
IBM	Global systems integrator	sells Generative AI; competes in Cybersecurity	structural	—
Infosys	Global systems integrator	sells Generative AI	structural	—
Tata Consultancy Services	Global systems integrator	sells Generative AI	structural	—
Salesforce	Customer-experience platform	sells Generative AI	structural	—

evidenced means this topic's own sources name the competitor — the signal ids are in the last column and are dated and clickable in the radar. **structural** means the curated register says they sell this technology into this vertical, which is true but is not proof they are in this deal. Register version competitors-2026-08-a.

Qualifying questions for the first meeting

1. What generative AI use cases are you currently deploying or piloting in clinical or operational settings?
2. Have you experienced any incidents where AI outputs have led to biased or inequitable care decisions?
3. Do you have an existing AI governance committee or framework? If so, how does it address bias and equity?
4. What regulatory or accreditation pressures are you facing regarding AI transparency and fairness?
5. How do you currently handle data sovereignty and security for AI models that process patient data?
6. Who in your organization would own the responsibility for ensuring AI does not discriminate against vulnerable groups?

Objections you will meet

They will say	You can answer
We already have an AI governance framework in place.	That's good, but many existing frameworks focus on technical robustness and privacy, not on decolonized ethics and equity. Our approach specifically addresses systemic biases that can affect vulnerable groups, which is increasingly a focus of regulators and payers.
This sounds like a research project, not something we can implement today.	You're right that the decolonized ethics concept is emerging, but we can start with a practical gap analysis of your current AI deployments and build a governance layer that is actionable now, using our Live Intelligence Studio and certified infrastructure.
We are already working with a major cloud provider on AI.	That's fine—we can complement that. Our value is in the governance and assurance layer, and we can work with your existing cloud investments, while adding sovereign, healthcare-specific certifications like HDS and SecNumCloud that are critical for patient data.

Risks and unknowns

The main risk is that the decolonized ethics framework is still a research concept, and there is no established market standard or regulatory requirement yet, so customers may not see it as a priority. Additionally, the framework's effectiveness in preventing discrimination is unproven in practice, and there is a lack of concrete metrics to measure success. The engagement could stall if the customer's AI deployments are too immature or if they already have internal governance teams that resist external frameworks.

Next action, by role (FR-17)

Role	Do this next
Presales / Proposal	Prepare a bid brief for a healthcare AI governance engagement, anchoring on Orange's HDS certification and ISO 27001, and referencing MicroPort CardioFlow as a proof point for responsible AI in medical devices, while outlining how Live Intelligence Studio can support transparent model auditing.
Sales	In a conversation with a healthcare provider's chief digital officer, open with: 'We see generative AI's promise in care, but it can silently encode bias against vulnerable groups. Orange is investing in governance frameworks that prevent that—can we explore how your organization is approaching AI ethics today?'
Strategist / Innovator	Commission a deep-dive brief on decolonized ethics and governance frameworks for generative AI in healthcare, leveraging Orange Group researchers and healthcare experts to map existing standards and identify gaps for vulnerable-group protection.

Evidence

Every claim in this brief resolves to one of these. Tier 1 is an authoritative primary source; tier 4 is vendor-published material, discounted in scoring (§4.3.7).

Signal	Date	Publisher	T	Title and link
SIG-A4D6D84811FE	2026-01-01	cordis.europa.eu	1	Building a Decolonized Ethics and Governance Framework for Just and Trustworthy AI in Healthcare
SIG-AECDD651DA98	2025-12-17	openalex.org	1	Edge-AI integrated secure wireless IoT architecture for real time healthcare monitoring and fed...
SIG-0A19EC698D37	2025-11-14	openalex.org	1	Agentic AI: a comprehensive survey of architectures, applications, and future directions
SIG-71EE3B7DB4F6	2025-11-03	openalex.org	1	Smart healthcare: the role of AI, robotics, and NLP in advancing telemedicine and remote patien...
SIG-702E3F720114	2025-11-01	cordis.europa.eu	1	LEADING CYBERSECURITY IN THE GENERATIVE AI ERA
SIG-ADF146AE36A5	2025-09-23	openalex.org	1	Artificial intelligence in healthcare and medicine: clinical applications, therapeutic advances...
SIG-F7E066A1966D	2025-09-03	openalex.org	1	A Framework for Lightweight Generative AI: Enabling Secure, Scalable, and Cloud-to-Edge Intelli...
SIG-373E00524578	2025-08-29	openalex.org	1	Personalized health monitoring using explainable AI: bridging trust in predictive healthcare
SIG-B7834050E076	2025-08-26	openalex.org	1	A Case Study of Environmental Footprints for Generative AI Inference: Cloud versus Edge
SIG-828846E90FC6	2025-07-01	cordis.europa.eu	1	TRUstworthy huMAN-centric artificial intelligence
SIG-F19E5C4832C5	2025-06-01	cordis.europa.eu	1	DeepKeep Safeguards AI Applications Across LLM, Vision, Spatial Sensing, Human-Machine Interact...
SIG-993FDF4B48F0	2026-08-06	arxiv.org	3	When Agentic AI Meets Integrated Sensing and Communication
SIG-F2B27060ED4C	2026-08-06	arxiv.org	3	From Siloed Algorithms to Compliance-First Agentic Platforms: A Multi-Layered Architecture for ...

How this brief was made

Computed, not written. Attractiveness, right to win, competitive intensity and every figure in the market-opportunity section are computed from stored inputs and can be re-derived (DR-05, NFR-01). No language model produced a number anywhere in this document (§4.4.4).

Written, then checked. The narrative sections and the solution diagram were generated from this topic's evidence and from the named asset and competitor lists. Factual sections must cite signal ids attached to this topic; sections that failed that check, or that contained a generated quantity or an unsupplied organisation, were removed rather than rewritten.

Removed by those checks in this brief: diagram (3 box(es) claimed to be an Orange asset that was not supplied — shown as third-party instead)

What	Version	When
Scores — weight set	w-2026-08-a	2026-08-18T19:08:26+00:00
Market sizing — assumptions	size-2026-08-a	2026-08-18T21:06:13+00:00
Competitor register	competitors-2026-08-a	2026-08-18T21:06:17+00:00
Narrative — prompt / model	describe-v1 / deepseek-chat	2026-08-18T20:42:05+00:00
Pipeline	0.1.0	2026-08-18

Generated 2026-08-19 04:58 UTC. Scores computed under a different weight set are not comparable with these (SC-10), and neither are market sizes computed under a different sizing version.