

Tamper-evident runtime evidence for AI agents in government using open-source Halo to ensure auditability and compliance.

Public / Government sector × AI governance and assurance × Blockchain and DLT · OS140 v1 · refreshed 2026-08-18

ATTRACTIVENESS	RIGHT TO WIN	COMPETITION	SERVICEABLE MARKET	HORIZON	PORTFOLIO
69 is the world moving	43 can we play, can we win	MEDIUM 1 named in the evidence	€220m observed evidence · per year	NEXT pilots published no volume pr...	L1 12 signals

Attractiveness and right to win are scored and displayed separately and are never combined (SC-12). Competition is a fourth quantity, not a discount on either.

Market opportunity

Method	Total (TAM)	Serviceable (SAM)	Obtainable (SOM)	Confidence
Observed: tendered contract value, annualised	€220m €220m – €482m	€220m €220m – €482m	€2.6m €2.6m – €5.8m	observed

All figures are annual, in EUR. Ranges compound the uncertainty of each factor below; the base case is the middle column of each.


Observed: tendered contract value, annualised

Factor	Value	Basis	Source	As of
Tendered contract value in matching CPV categories	€220m per year	observed	TED (Tenders Electronic Daily) · ted	2026-05-19..2026-08-05
Assumed obtainable share of the serviceable market	1.2 % of SAM	assumption	config/sizing.yaml · obtainable_share	size-2026-08-a

Read this before quoting the number. Observed over 85 days of TED coverage and scaled to a year (x4.3); a short collection window makes that scaling rough. 34 matching notices, matched on vertical via the CPV crosswalk — \$4.5.2's warning applies: a crosswalk error lands directly in this number. Public procurement only. Private-sector demand for the same capability is not visible here at all, so this is a floor rather than a market size.

The solution, in one picture

Tamper-Evident AI Governance for Government



Orange asset Partner Customer-owned Third party / to be sourced

This solution provides a tamper-evident audit trail for AI agents, running on sovereign, certified infrastructure, to meet government compliance requirements.

Boxes marked as Orange assets are validated against the linked business graph; a box the model claimed for Orange without a matching asset is shown as third party.

Who buys, and why now

The buyer is typically a Chief Information Officer or Chief Digital Officer in a government agency, while the pain is felt by compliance officers and auditors who must demonstrate that AI systems operate within legal and ethical boundaries. The trigger is often an audit, a new regulation, or a high-profile incident involving an AI system that raises questions about accountability. They need to show that every action taken by an AI agent can be traced and verified, which is impossible with current logging approaches that can be altered or deleted.

Sources: SIG-A084505FAA05 github.com 2026-07-07

Why it is hot now — every claim bound to a dated source

- An open-source tool provides tamper-evident runtime evidence for AI agents, addressing the need for auditable AI operations. [SIG-A084505FAA05]
- An open-source tool provides tamper-evident runtime evidence for AI agents, addressing the need for audibility and governance of AI systems. [SIG-A084505FAA05]

What Orange would deliver, and why it can win

Named assets behind this — each individually inspectable (LK-08)

Asset	Type	Link	Why it is linked	Curator
Live Intelligence Studio	offer	L1	offer addresses the use case but not this technology	unconfirmed
ISO 27001	certification	SUP	certification Orange holds is required by this vertical	unconfirmed

SecNumCloud	certification	SUP	certification Orange holds is required by this vertical	unconfirmed
CISPE	certification	SUP	certification Orange holds is required by this vertical	unconfirmed
AI / Data / Cloud experts	capability pool	SUP	capability pool staffs this domain, vertical or technology	unconfirmed
Cybersecurity experts	capability pool	SUP	capability pool staffs this domain, vertical or technology	unconfirmed
Digital experts	capability pool	SUP	capability pool staffs this domain, vertical or technology	unconfirmed
Defense and security experts	capability pool	SUP	capability pool staffs this domain, vertical or technology	unconfirmed
Healthcare experts	capability pool	SUP	capability pool staffs this domain, vertical or technology	unconfirmed
Orange Group researchers	capability pool	SUP	capability pool staffs this domain, vertical or technology	unconfirmed
Grand Paris Sud	reference	SUP	published reference in this vertical, different use case	unconfirmed

Competitive landscape

MEDIUM competitive intensity (score 3.2; 1 of 8 listed players appear in this topic's own evidence)

A real field, including at least one player the customer will already know. Expect to be compared.

Competitors include global systems integrators like Accenture, Atos/Eviden, Capgemini, IBM, and Sopra Steria, all active in the public sector and offering cloud and cybersecurity services. Telecom peers Deutsche Telekom/T-Systems and Telefonica Tech also compete in this space. AWS, while a partner, is also a competitor in cloud services and has been highlighted for its focus on agentic AI. These competitors are strong in large-scale system integration and have existing government contracts, but they may not offer the same level of sovereign, certified infrastructure that Orange can provide.

Sources: SIG-8A29805612CD bing.com 2026-06-23

Competitor	Category	Position here	Basis	Evidence
AWS	Hyperscaler	competes in Cloud — also an Orange partner	evidenced	SIG-8A29805612CD
Accenture	Global systems integrator	active in Public / Government sector; competes in Cloud	structural	—
Atos / Eviden	Global systems integrator	active in Public / Government sector; competes in Cloud, Cybersecurity	structural	—
Capgemini	Global systems integrator	active in Public / Government sector; competes in Cloud	structural	—
Deutsche Telekom / T-Systems	Telecom operator peer	active in Public / Government sector; competes in Cloud, Cybersecurity	structural	—
IBM	Global systems integrator	active in Public / Government sector; competes in Cloud, Cybersecurity	structural	—
Sopra Steria	Global systems integrator	active in Public / Government sector; competes in Cloud, Cybersecurity	structural	—
Telefonica Tech	Telecom operator peer	active in Public / Government sector; competes in Cloud, Cybersecurity	structural	—

evidenced means this topic's own sources name the competitor — the signal ids are in the last column and are dated and clickable in the radar. **structural** means the curated register says they sell this technology into this vertical, which is true but is not proof they are in this deal. Register version competitors-2026-08-a.

Qualifying questions for the first meeting

- 1 What is your current process for auditing the decisions made by AI agents in your agency?
- 2 Have you experienced any compliance or audit findings related to AI systems in the past year?
- 3 Do you have a requirement to provide tamper-evident logs of AI agent actions to regulators or oversight bodies?
- 4 What AI frameworks or platforms are you currently using for agentic AI, and how do you monitor their runtime behavior?

- 5
- Are you planning to deploy AI agents in areas that require high assurance, such as citizen-facing services or national security?
- 6
- What is your timeline for implementing AI governance measures, and is there a specific regulatory deadline driving it?

Objections you will meet

They will say	You can answer
We already have logging and monitoring tools for our AI systems; why do we need another one?	That's true, but standard logging is not tamper-evident. Halo provides cryptographic evidence that can prove the integrity of the logs, which is essential for regulatory compliance and auditability.
Open-source tools are not secure enough for government use.	Open-source does not mean insecure. Halo's code is publicly auditable, and we will run it on our SecNumCloud-certified infrastructure, adding an extra layer of security and compliance.
We don't have the budget for another AI governance project right now.	We understand budget constraints. This solution can be piloted with a small set of AI agents to demonstrate value and compliance benefits, and the open-source nature of Halo reduces licensing costs.

Next action, by role (FR-17)

Role	Do this next
Presales / Proposal	Bid brief: prepare a differentiated proposal for a public-sector AI governance tender by combining Live Intelligence Studio with our ISO 27001, SecNumCloud, and CISPE certifications, and reference the Grand Paris Sud project as proof of our public-sector delivery capability.
Sales	In a meeting with a public-sector CIO or CISO, say: 'We are exploring how to make AI agent operations auditable and tamper-evident, building on our Live Intelligence Studio and our experience with Grand Paris Sud, and we would like to understand your current challenges with AI governance and compliance.'
Strategist / Innovator	Deep-dive brief: assess the feasibility of building a tamper-evident runtime evidence capability for AI agents on top of Live Intelligence Studio, leveraging our AI / Data / Cloud, Cybersecurity, and Digital experts, and validate the market pull with government pilots before investing in a full offer.

Evidence

Every claim in this brief resolves to one of these. Tier 1 is an authoritative primary source; tier 4 is vendor-published material, discounted in scoring (§4.3.7).

Signal	Date	Publisher	T	Title and link
SIG-5280658BD549	2026-07-08	Computers	1	Post-Quantum Cryptography Migration for Agentic AI Systems
SIG-81E56EEBF9D4	2026-06-15	openalex.org	1	Post-Quantum Cryptography Migration for Agentic AI Systems
SIG-0A19EC698D37	2025-11-14	openalex.org	1	Agentic AI: a comprehensive survey of architectures, applications, and future directions
SIG-D4C81A60CDDC	2025-10-17	openalex.org	1	An Energy-Aware Generative AI Edge Inference Framework for Low-Power IoT Devices
SIG-B7834050E076	2025-08-26	openalex.org	1	A Case Study of Environmental Footprints for Generative AI Inference: Cloud versus Edge
SIG-4674EE3A0873	2025-02-26	The Role of Network Secur...	1	5G-Enabled Smart Healthcare System with the Integration of Blockchain Technology
SIG-01140D9380B3	2025-01-01	cordis.europa.eu	1	Human-Centric Artificial Intelligence for Sustainable Future
SIG-8A29805612CD	2026-06-23	bing.com	2	Agentic AI Takes Centre Stage at AWS Summit 2026
SIG-ED4013FC65F2	2026-06-18	Federal News Network	2	Cloud Exchange 2026: Ping Identity's Kelvin Brewer on identity as foundation of secure AI adopt...
SIG-E339EBF1D5CC	2026-05-11	Future Market Insights	2	AI In Government And Public Services Market - Future Market Insights
SIG-DDF1F9E38278	2026-08-01	arxiv.org	3	Latency-Tolerant Cloud-Edge Collaborative Vision-Language-Action Models via Emergent Representa...
SIG-A084505FAA05	2026-07-07	github.com	3	Show HN: Halo — open-source, tamper-evident runtime evidence for AI agents

How this brief was made

Computed, not written. Attractiveness, right to win, competitive intensity and every figure in the market-opportunity section are computed from stored inputs and can be re-derived (DR-05, NFR-01). No language model produced a number anywhere in this document (§4.4.4).

Written, then checked. The narrative sections and the solution diagram were generated from this topic's evidence and from the named asset and competitor lists. Factual sections must cite signal ids attached to this topic; sections that failed that check, or that contained a generated quantity or an unsupplied organisation, were removed rather than rewritten.

Removed by those checks in this brief: summary (named entities that were never supplied: Halo); what_is_changing (named entities that were never supplied: Halo); what_orange_would_deliver (named entities that were never supplied: Halo); why_orange_can_win (named entities that were never supplied: Halo); risks_and_unknowns (named entities that were never supplied: Halo)

What	Version	When
Scores — weight set	w-2026-08-a	2026-08-18T19:08:26+00:00
Market sizing — assumptions	size-2026-08-a	2026-08-18T21:06:16+00:00
Competitor register	competitors-2026-08-a	2026-08-18T21:06:19+00:00
Narrative — prompt / model	describe-v1 / deepseek-chat	2026-08-19T04:57:24+00:00
Pipeline	0.1.0	2026-08-18

Generated 2026-08-19 04:58 UTC. Scores computed under a different weight set are not comparable with these (SC-10), and neither are market sizes computed under a different sizing version.